


Neural Networks

Linear Classification and Optimization



Objective

We want to classify cats and as well as other animals.

We have N training examples, K classes (Number of animals), D pixels (dimension of our picture)

$N : 10,000$

$K : 10$

$D : m \times 64 \times 64 \times 3$

Our objective is map N animals to its specific class

How we can do it?

Simplest linear function => $f(x) = Wx + b$

Adding more layers to the networks makes the network recognize patterns and different features

It makes feature engineering for us

How we can do it?

$$f(x) = W_n * (W_{n-1} * (...(W_1 * x + b_1)...) + b_{n-1}) + b_n$$



Feature? Show me!



Low level features


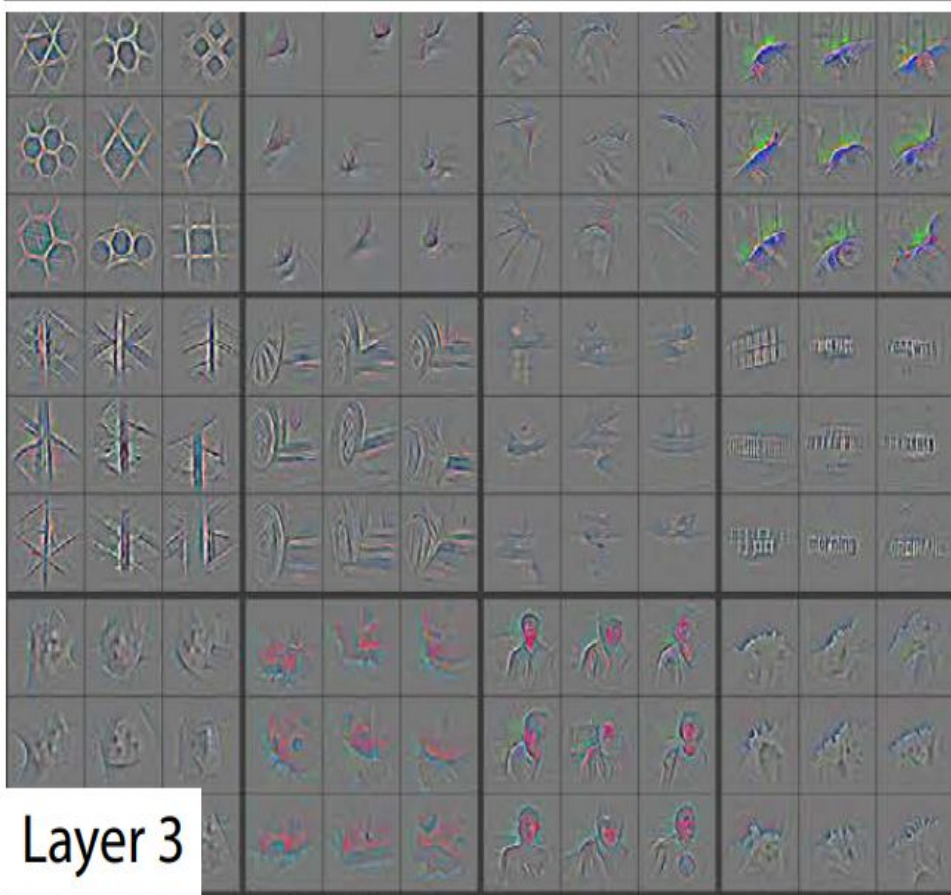
Edges, color changes, diagonals, etc.




Layer 1

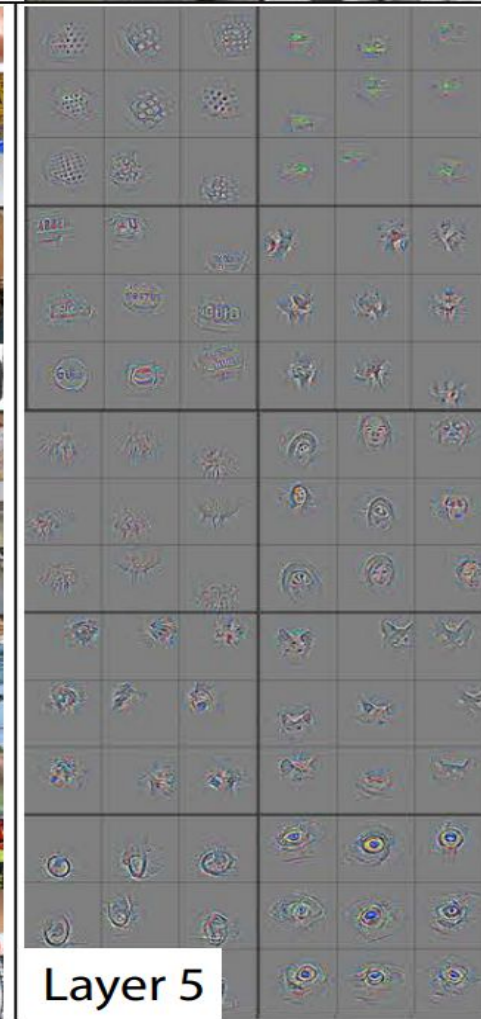










Layer 4




Layer 5




Activation Functions

$$A = \frac{1}{1+e^{-x}}$$




Sigmoid Function

$$f(x) = \tanh(x) = \frac{2}{1+e^{-2x}} - 1$$




tanh Function

Activation Functions



ReLU Function



Leaky ReLU Functions

Activation Functions

Helps the function to make itself nonlinear

Since the function we use is the simple linear equation, we need to change it and make it non-linear prediction function.

Other Activation Functions: https://en.wikipedia.org/wiki/Activation_function

Loss Function


Binary loss function was to classify only two different classes (dog, cat), (malignant, safe)


Since we have $K \geq 2$, we need to use something different

Softmax Classifier

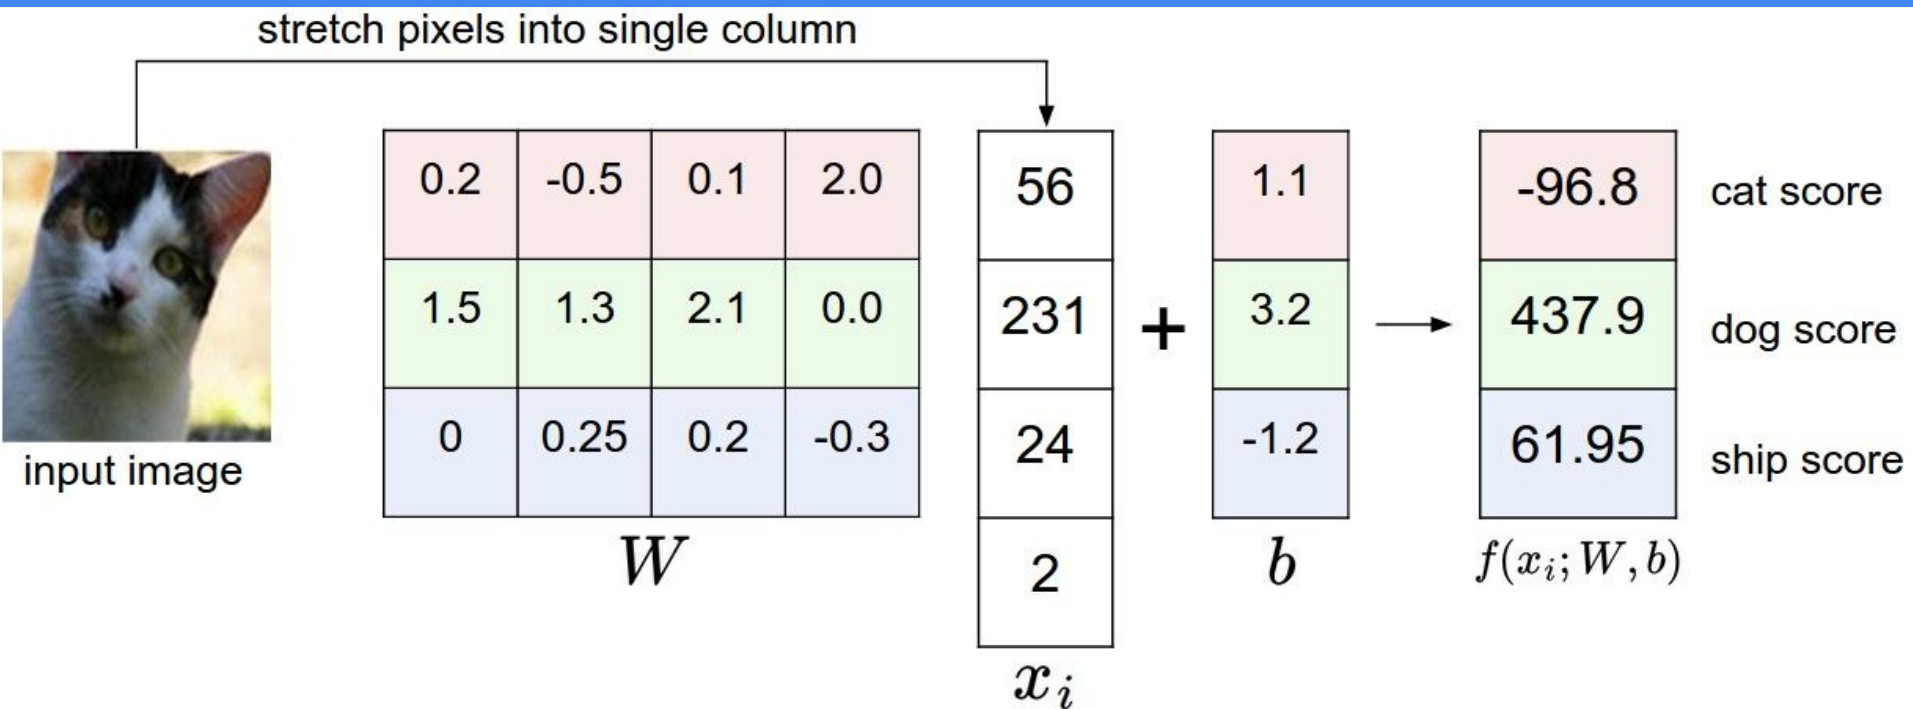
Meet Softmax

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$





Loss function landscape for the Multiclass SVM (without regularization) for one single example (left,middle) and for a hundred examples (right) in CIFAR-10. Left: one-dimensional loss by only varying \mathbf{a} . Middle, Right: two-dimensional loss slice, Blue = low loss, Red = high loss. Notice the piecewise-linear structure of the loss function. The losses for multiple examples are combined with average, so the bowl shape on the right is the average of many piece-wise linear bowls (such as the one in the middle).



matrix multiply + bias offset

0.01	-0.05	0.1	0.05
0.7	0.2	0.05	0.16
0.0	-0.45	-0.2	0.03

W

-15
22
-44
56

x_i

+

0.0
0.2
-0.3

b

y_i

2

hinge loss (SVM)

-2.85
0.86
0.28

$$\begin{aligned} & \max(0, -2.85 - 0.28 + 1) + \\ & \max(0, 0.86 - 0.28 + 1) \\ & = \\ & \mathbf{1.58} \end{aligned}$$

cross-entropy loss (Softmax)

-2.85
0.86
0.28

\exp

0.058
2.36
1.32

normalize
(to sum to one)

0.016
0.631
0.353

$$\begin{aligned} & -\log(0.353) \\ & = \\ & \mathbf{1.04} \end{aligned}$$